

Project 5: Sliding Tile Puzzle Solver

In this assignment you will implement one of two data structures that enable a state-space search algorithm called A* (pronounced A-star) that is applied to solve a simple sliding-tile puzzle. The algorithm and puzzle solving code are provided, you only have to implement and test the data structure.

Sliding Tile Puzzles:

[Sliding-tile puzzles](#) are classical problems in state-space search. We will be considering the eight-tile puzzle. It consists of an 3x3 board of eight tiles labeled A-H and a single blank (missing) tile. This is called the *state* of the puzzle. A tile may slide into the blank position vertically or horizontally giving rise to other states. The goal of the puzzle is to find the tile moves need to reach a particular goal state. For example, consider the following puzzle state

```
ABC
DEF
GH
```

There are two possible states resulting from sliding H right into the blank space or sliding F down into the blank space.

```
ABC  ABC
DEF  DE
G H  GHF
```

The state space search proceeds by searching through moves until a goal state is reached. This is similar to the maze problem from P4 but the number of possible board states is much larger. Thus, we need a more efficient search algorithm that can use additional information, called a *heuristic*. The optimal algorithm (using the same information) is called A*.

A* algorithm

A* algorithm

Like the breadth-first search from P4 the A* algorithm can be described generically using a type State and operations on a problem instance:

- `problem.initial()` returns the initial state of the problem
- `problem.goal(state)` returns true if state is the goal state, else false
- `problem.actions(state)` returns a list of states resulting from possible transitions from state

We only need add two additional variables, traditionally called the path-cost and f-cost. The path cost, g , is the number of state-transitions from the initial state to the current state. The f-cost is the path-cost plus a state-dependent value called the heuristic, h , that estimates how far from the goal state the current state is; that is $f = g+h$.

function astar_search(problem) returns a solution or failure

 s = problem.initial()

 if problem.goal(s) return s

 frontier is a min priority queue with s as the initial element

 explores is an empty set

 while true

 if frontier is empty return failure

 s = pop next state from frontier

 add s to explored

 for each state s' in problem.actions(s) do

 if s' not in explored or frontier then

 if problem.goal(s') then return s'

 insert s' into the frontier

 else if s' is in the frontier with a higher path-cost

 replace the state in the frontier with the current s'

The supporting data structures for A* are a priority queue modified to allow for inclusion tests and replacement, and a set. Unlike in P4 it is infeasible to implement the inclusion tests using a simple test array because again, the number of states in general is very large.

We can apply this algorithm to our puzzle solver easily. The puzzle state is simply a given arrangement of tiles. State transitions (actions) are determined by the location of the blank slot. We will simplify things somewhat by only keeping track of the path-cost rather than the actual sequence of moves.

A* Implementation

An implementation of the A* algorithm is provided in the module PuzzleSolver (puzzle_solver.hpp and puzzle_solver.cpp) using the State module (the template in state.hpp and state.cpp). The puzzle board and supporting functionality is provided in the Puzzle module (puzzle.hpp and puzzle.cpp).

The algorithm requires implementations of the data structures frontier_queue and explored_set. The explored_set in the provided code uses [unordered_set](#), the hash-table implementation in the

standard library. Since the frontier is not a normal priority queue, this is the data structure you will be implementing. The API of this module is defined in `frontier_queue.hpp` and described below.

There is a set of tests for the puzzle solver in `test_solver.cpp`. **When your frontier queue is implemented properly, these tests should pass.** Each test takes two strings in the form "012345678", where 0 = A, 1 = B, ... 7 = H, and 8 = BLANK and converts them to Puzzle instances. It then checks the solution path-cost from one puzzle to another is the correct one, including the symmetric case (swap initial and goal puzzle).

Frontier Queue

The frontier queue template should be implemented in `frontier_queue.hpp` and `frontier_queue.cpp` as a min heap using dynamic allocation as necessary and have the following complexity for each member (see the header file comments for details):

- `push` should add a state to heap with complexity $O(\log n)$
- `pop` should remove and return the state in the heap with the lowest f-cost with complexity $O(\log n)$
- `contains` should return true if the state is in the frontier with complexity $O(n)$ or better
- `replaceif` should replace the given state in the heap if it has a higher path-cost with the resulting queue still being a valid heap. This should have a complexity of $O(n)$ or better

Note: You may use any combination of containers or algorithms from the C++11 standard library to implement the frontier queue. You will need to read and understand most of the code provided to you.